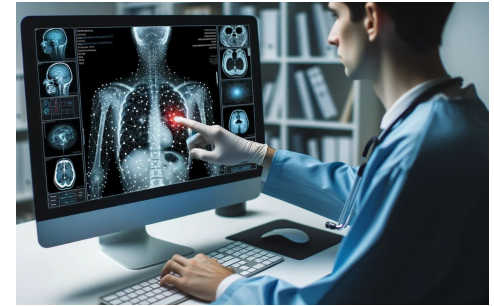
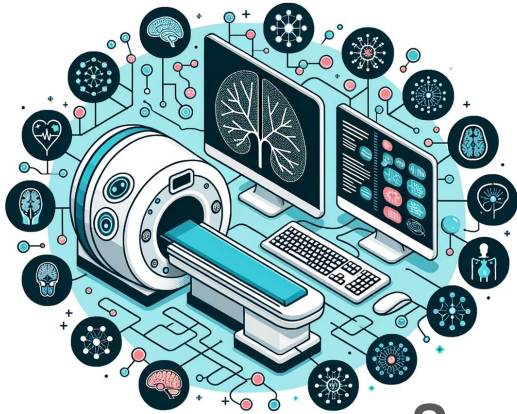


Applied Deep Learning and Generative Models in Healthcare



Session 2: Medical Image Segmentation
Date: Jan 18 2025

Instructor: Mahmoud E. Khani, Ph.D.

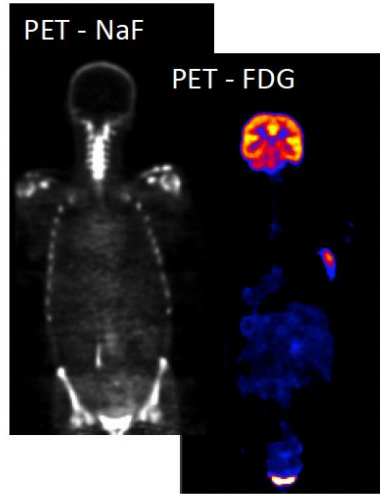
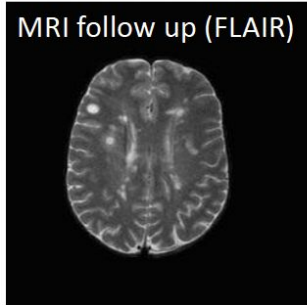
Motivation

- Images represent 90% of medical data

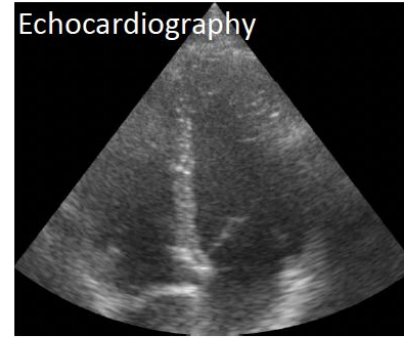


Medical images, many modalities

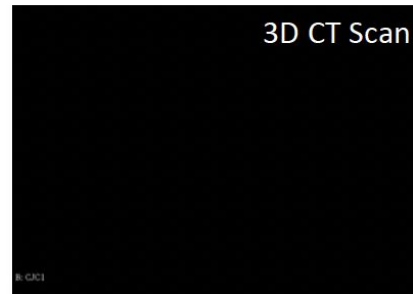
- 3D/2D (+t), different physics



Positron Emission Tomography

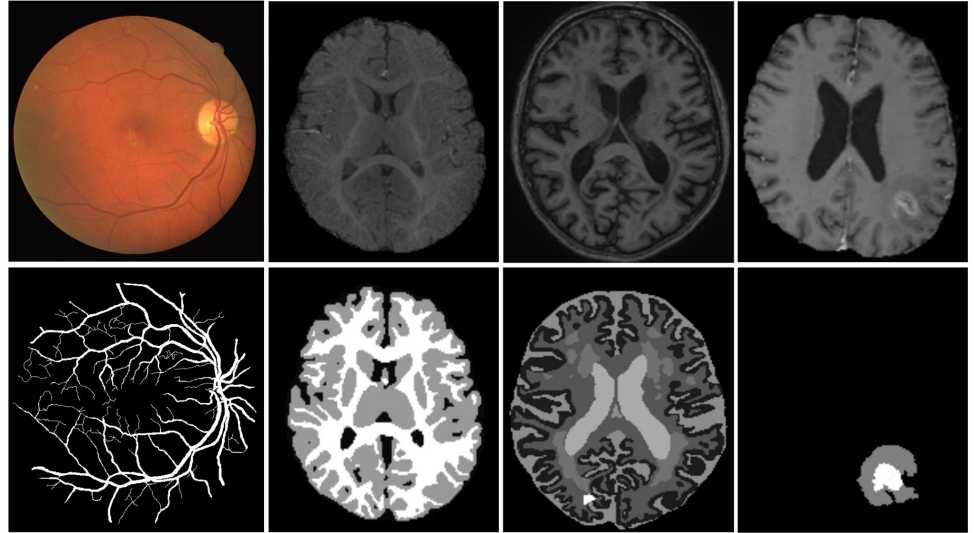
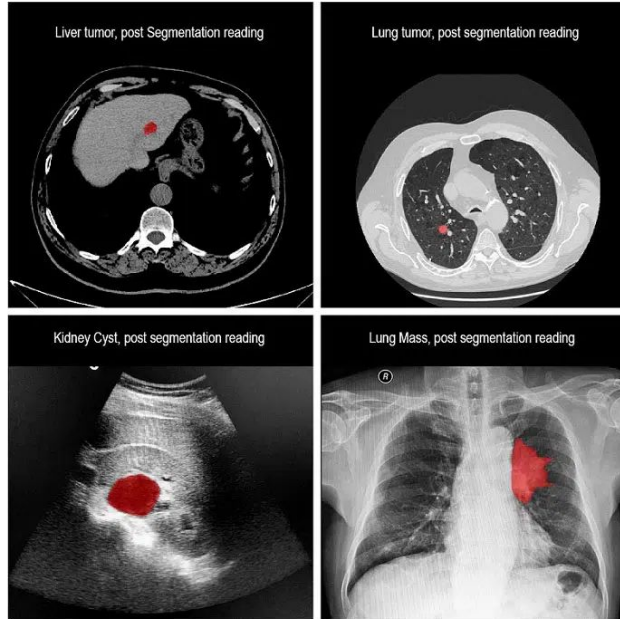


Fundus examination



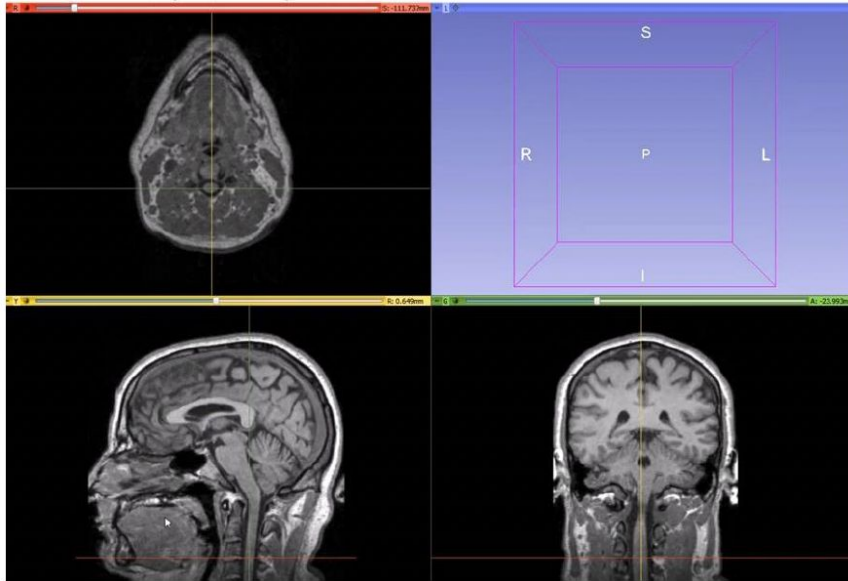
Medical Image Segmentation - Significance

- Most publications related to applications of deep learning and computer vision in medical imaging



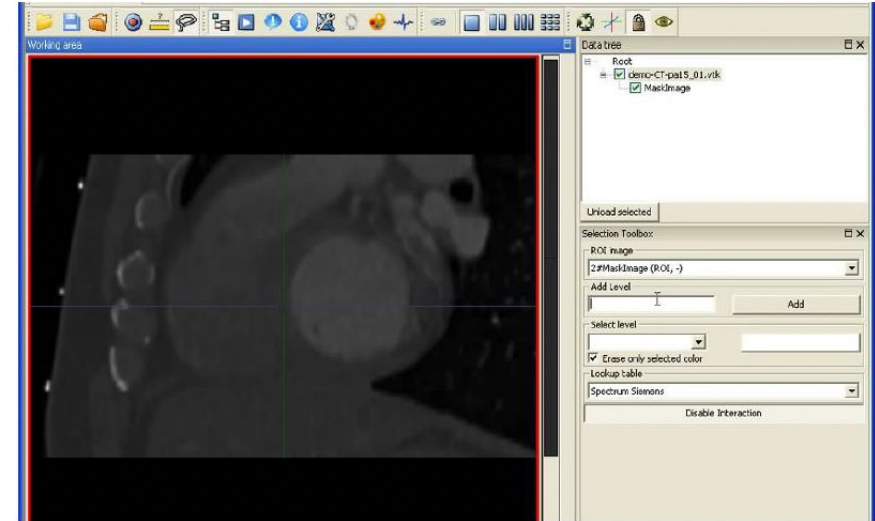
Manual image segmentation is time consuming!

3D Slicer - Quick Manual Segmentation with Interpolation



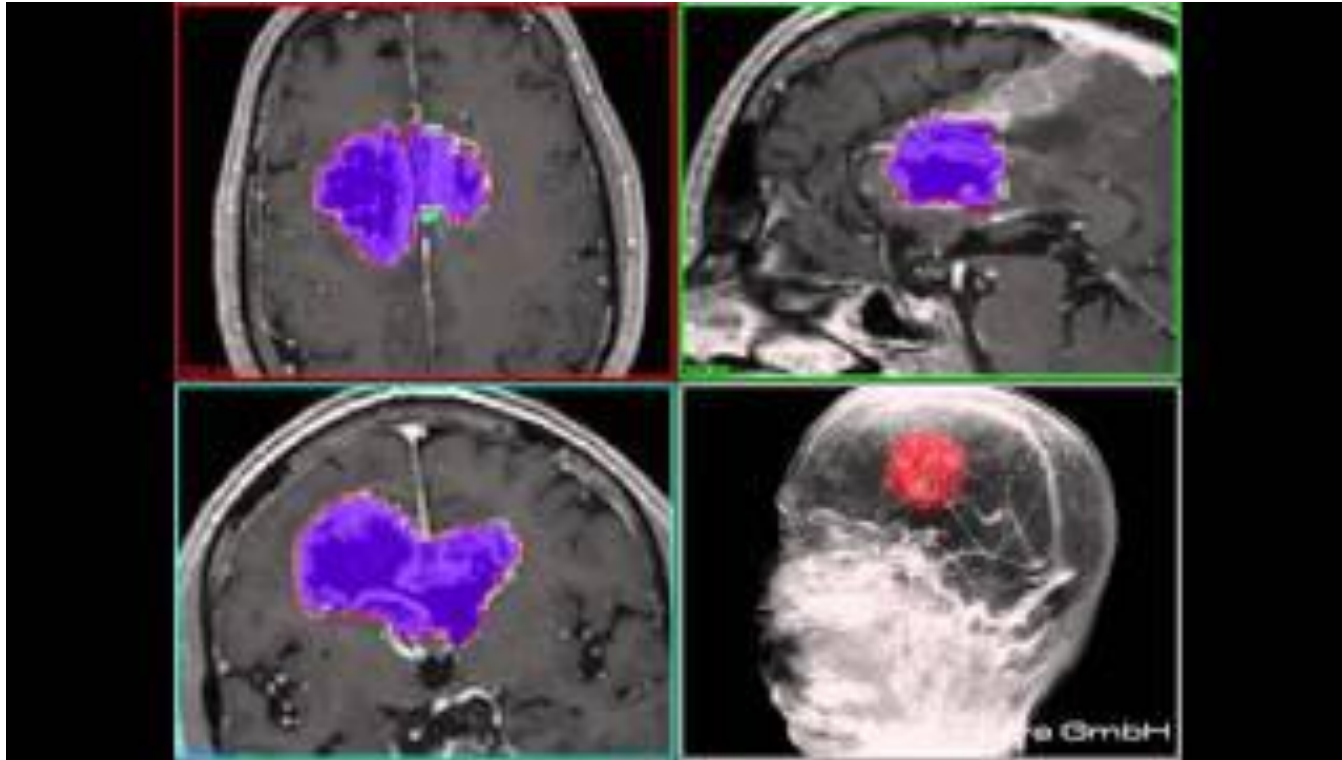
https://www.youtube.com/watch?v=u93kl1MG6lc&ab_channel=PerkLabResearch

GIMIAS Manual Segmentation



https://www.youtube.com/watch?v=rAHA1OZC8h8&ab_channel=GIMIAS

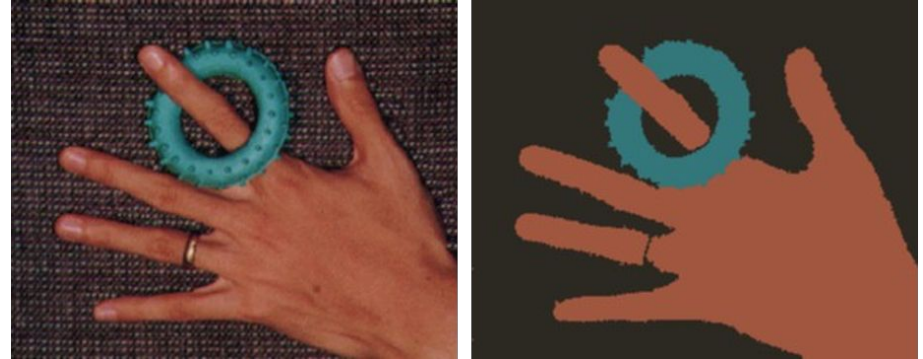
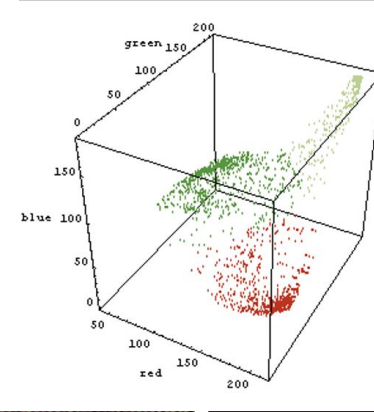
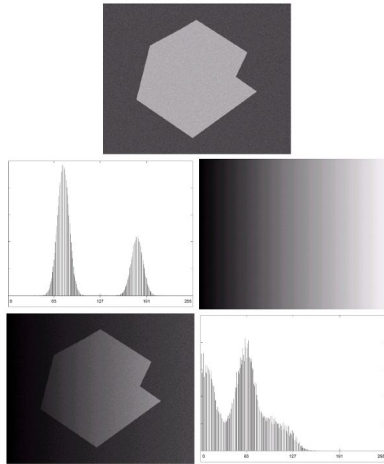
Semi-automated medical image segmentation



https://www.youtube.com/watch?v=7wCC2NaVLjs&ab_channel=ChimaeraGmbH

Classical image segmentation algorithms

- Thresholding
- Region-oriented segmentation
- Watershed segmentation
- K-means clustering



U-Net: Convolutional Networks for Biomedical Image Segmentation

Olaf Ronneberger, Philipp Fischer, and Thomas Brox

Computer Science Department and BIOSS Centre for Biological Signalling Studies
University of Freiburg, Germany
ronneber@informatik.uni-freiburg.de
<http://lmb.informatik.uni-freiburg.de/>

Abstract. There is large consent that successful training of deep networks requires many thousand annotated training samples. In this paper, we present a network and training strategy that relies on the strong use of **data augmentation** to use the available annotated samples more efficiently. The **architecture consists of a contracting path to capture context and a symmetric expanding path that enables precise localization**. We show that such a network can be trained end-to-end from very few images and outperforms the prior best method (a **sliding-window convolutional network**) on the **ISBI challenge for segmentation** of neuronal structures in electron microscopic stacks. Using the same network trained on transmitted light microscopy images (phase contrast and DIC) we won the ISBI cell tracking challenge 2015 in these categories by a large margin. Moreover, the network is fast. Segmentation of a 512x512 image takes less than a second on a recent GPU. The full implementation (based on Caffe) and the trained networks are available at <http://lmb.informatik.uni-freiburg.de/people/ronneber/u-net>.



IEEE International Symposium on Biomedical Imaging (ISBI). It focuses on benchmarking image segmentation methods applied to biomedical and medical imaging tasks. The challenge typically involves segmenting anatomical structures, pathological regions, or cellular/molecular components in medical images, such as MRI, CT, microscopy, or histopathology scans.

Medical image segmentation challenges

- Tumor segmentation
- Retinal vessel segmentation
- Organ delineation
- Cell tracking
- Disease detection.

$$\text{Dice} = \frac{2 \times \text{Area of overlap}}{\text{Total area}} = \frac{2 \times \text{Prediction} \cap \text{Ground truth}}{\text{Prediction} \cup \text{Ground truth}}$$

Evaluation metrics:

- Dice similarity coefficient (DSC)
- intersection-over-union (IoU)

```
def dice_coef(groundtruth_mask, pred_mask):  
    intersect = np.sum(pred_mask*groundtruth_mask)  
    total_sum = np.sum(pred_mask) + np.sum(groundtruth_mask)  
    dice = np.mean(2*intersect/total_sum)  
    return round(dice, 3) #round up to 3 decimal places
```

https://medium.com/@nghihuynh_37300/understanding-evaluation-metrics-in-medical-image-segmentation-d289a373a3f

Segmentation success

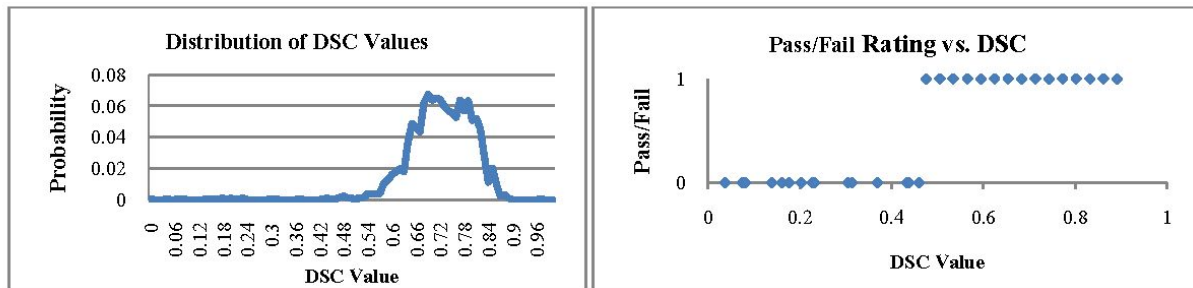


Fig. 1. (a) Distribution of DSC values from the 2,500 segmentations. (b) “Pass/fail” ratings assigned by the radiologists for the 30 randomly selected images used in the study. A value of one indicates a success and a value of zero indicates a failure.

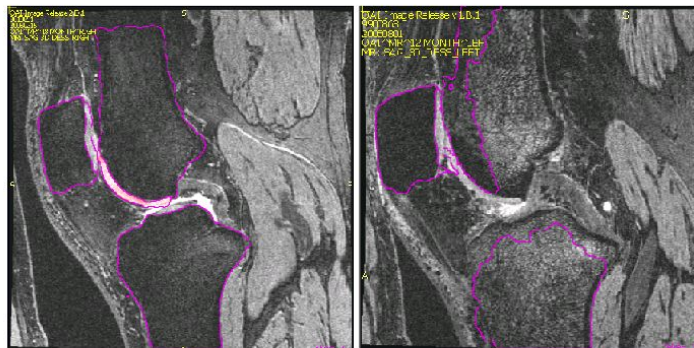
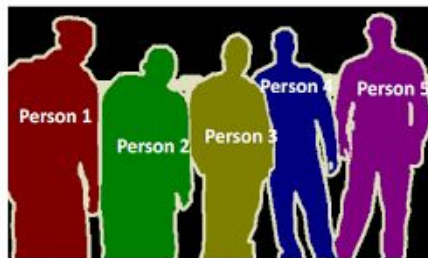


Fig. 2. (left) A typical segmentation labeled as passing by the expert radiologists. (right) A typical segmentation labeled as failing by the expert radiologists.

Semantic vs Instance Segmentation

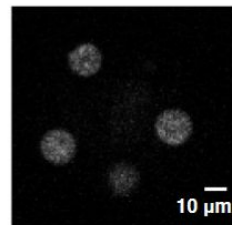


Semantic Segmentation

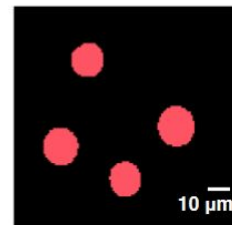


Instance Segmentation

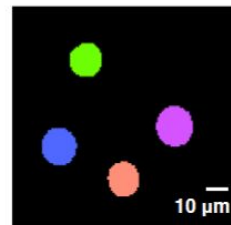
2D Fluorescence
Microscopic Image



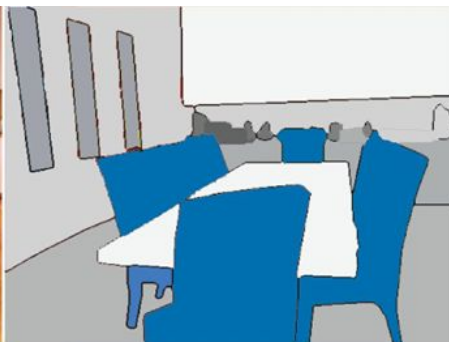
Semantic Segmentation



Instance Segmentation



Input Image

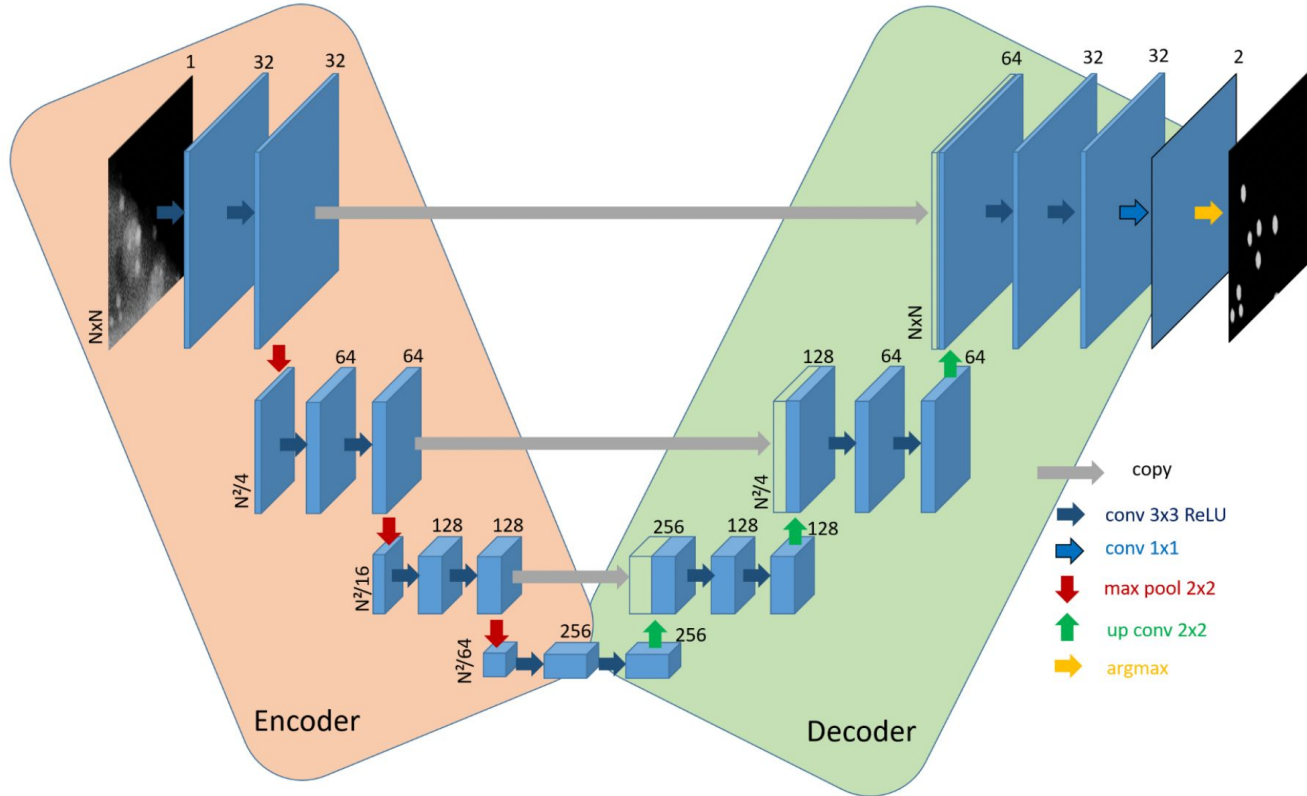


Semantic Segmentation

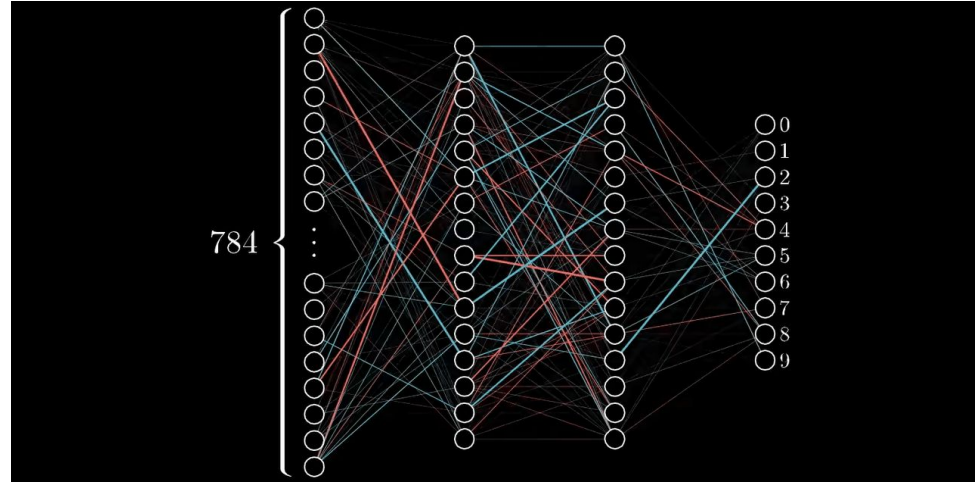
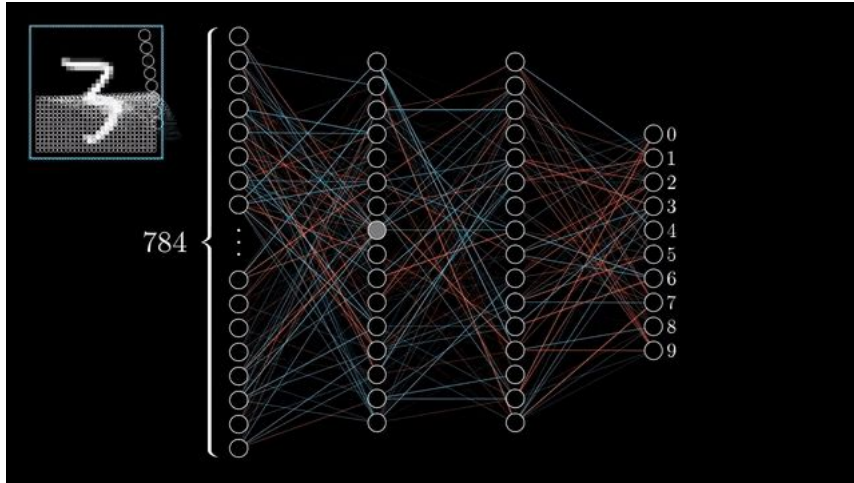


Instance Segmentation

U-net: Automatic image semantic segmentation

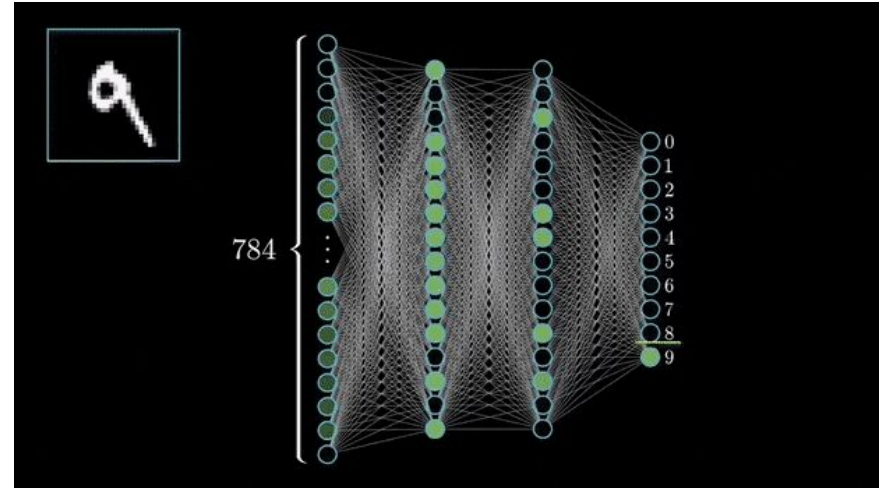
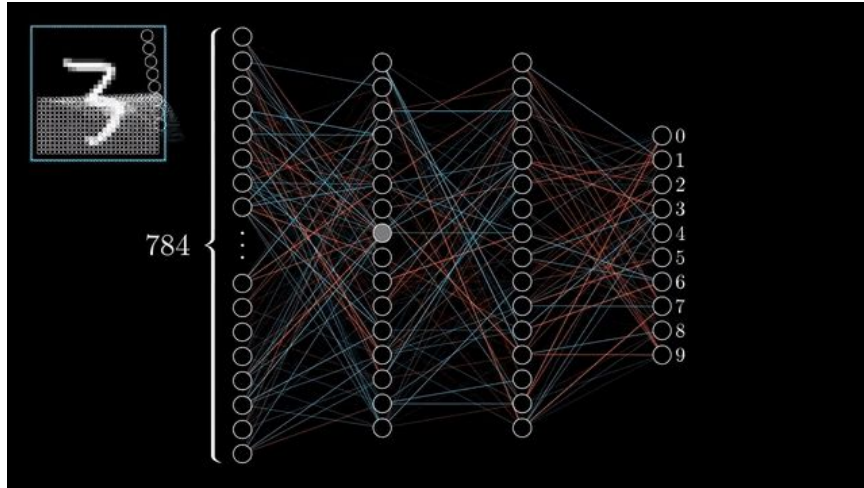


Multi-layer perceptrons (MLP) - Backpropagation



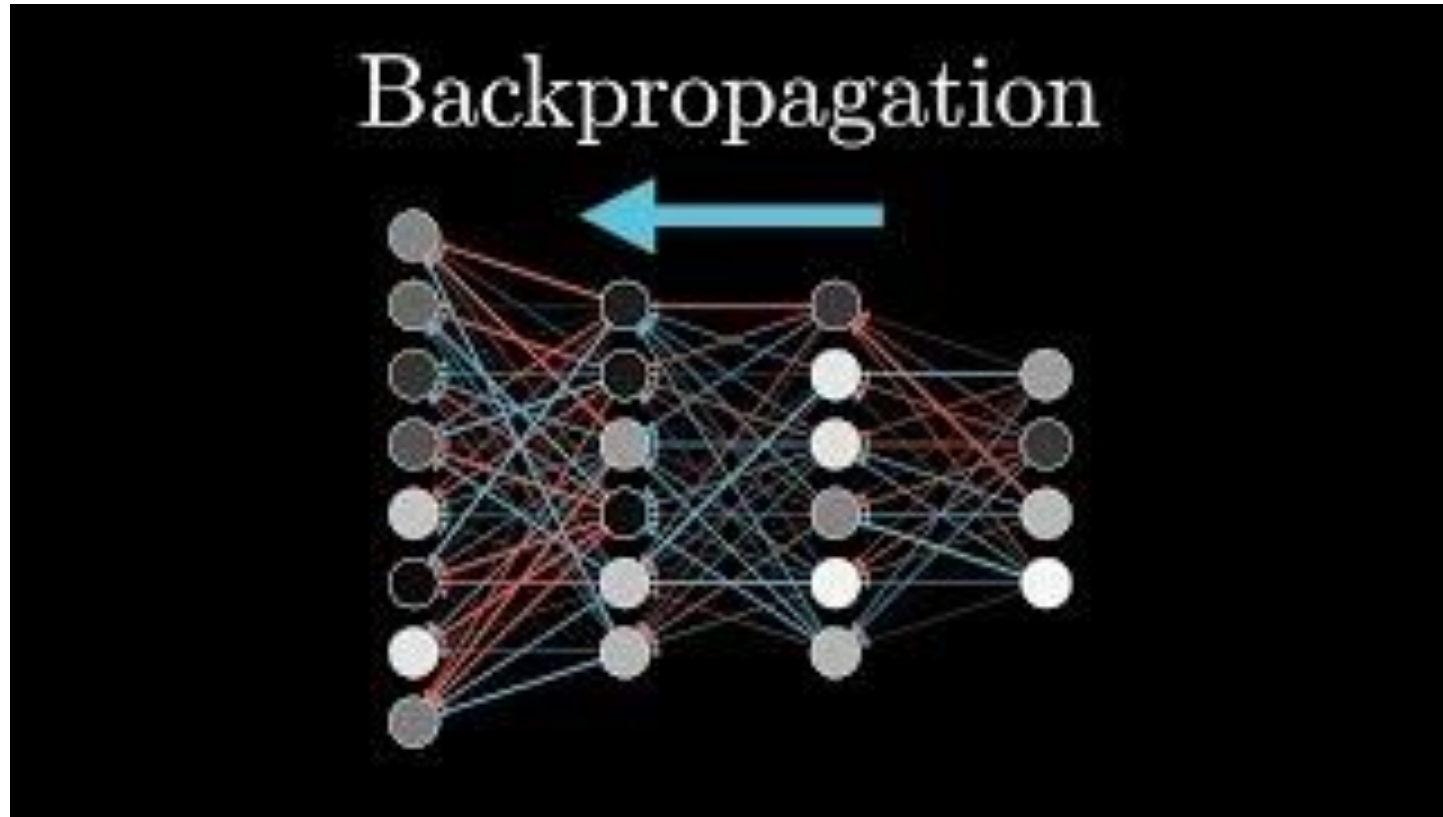
<https://pub.towardsai.net/backpropagation-2eeb25201095>

Multi-layer perceptrons - Backpropagation

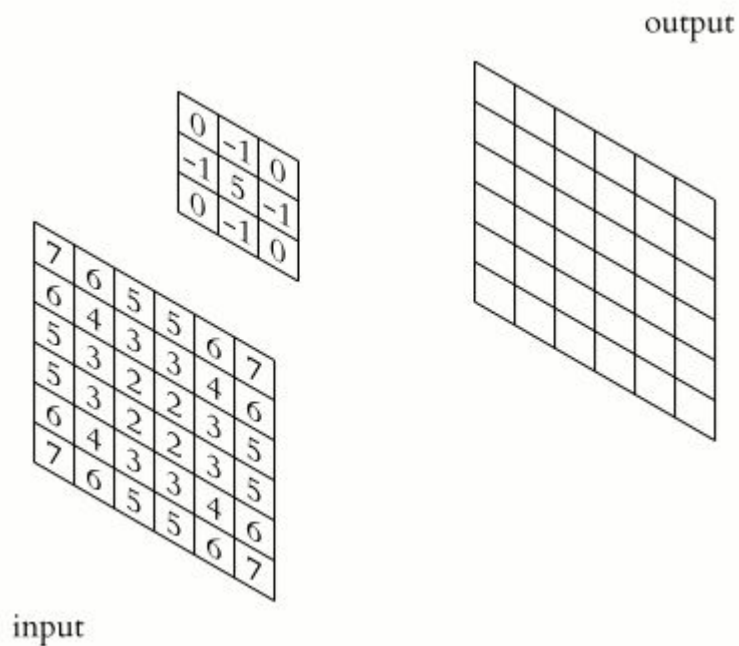


<https://medium.com/thedeephub/convolutional-neural-networks-a-comprehensive-guide-5cc0b5eae175>

Backpropagation explained intuitively



Convolutional Layers



Input Kernel Output

0	1	2
3	4	5
6	7	8

*

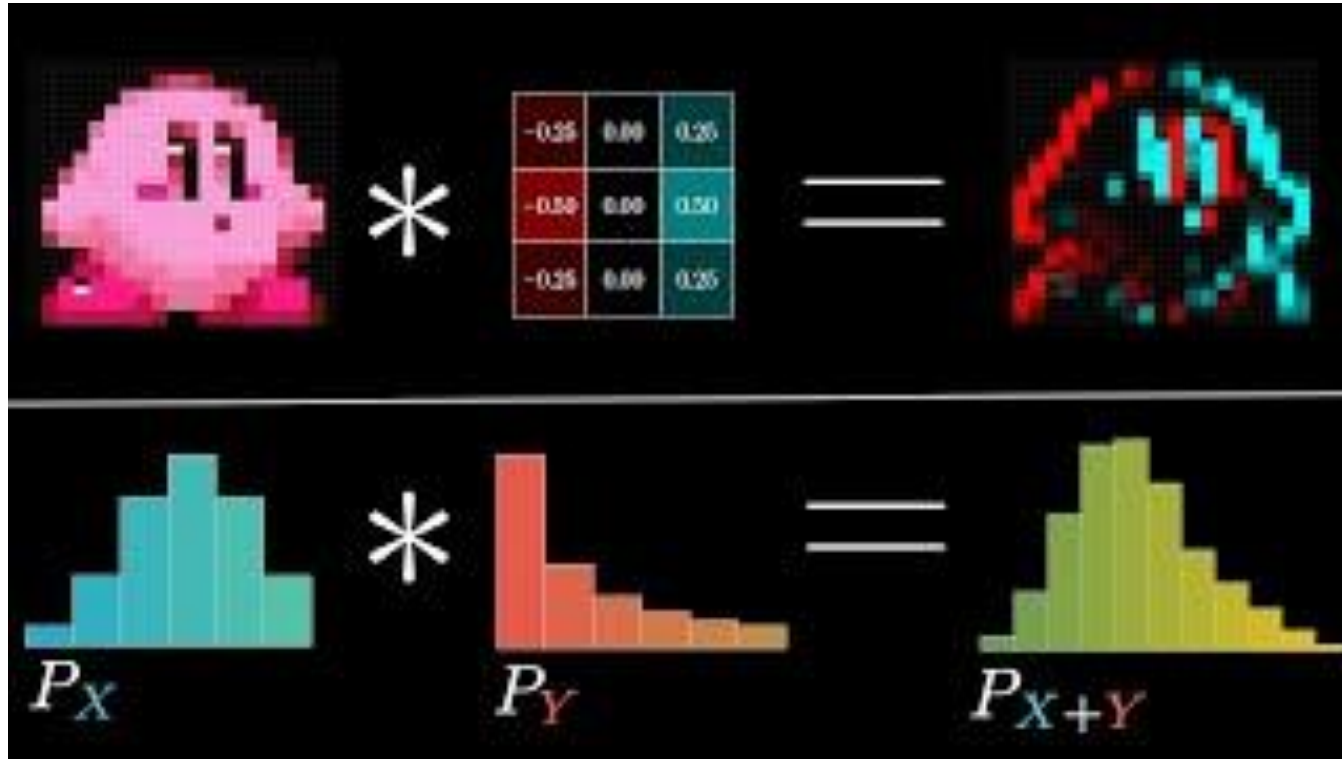
0	1
2	3

=

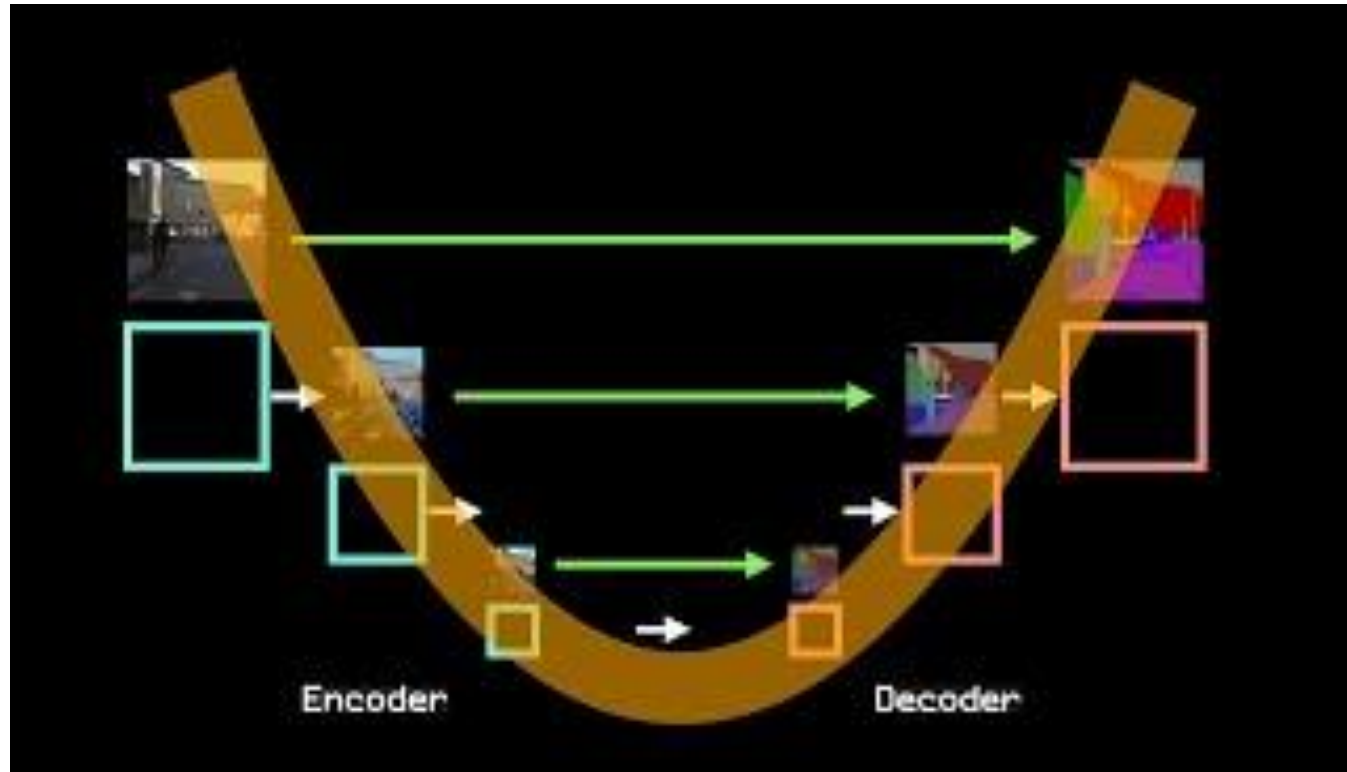
19	25
37	43

- $(0 \times 0) + (1 \times 1) + (3 \times 2) + (4 \times 3) = 19$
- $(1 \times 0) + (2 \times 1) + (4 \times 2) + (5 \times 3) = 25$
- $(3 \times 0) + (4 \times 1) + (6 \times 2) + (7 \times 3) = 37$
- $(4 \times 0) + (5 \times 1) + (7 \times 2) + (8 \times 3) = 43$

What is a convolution in image processing?



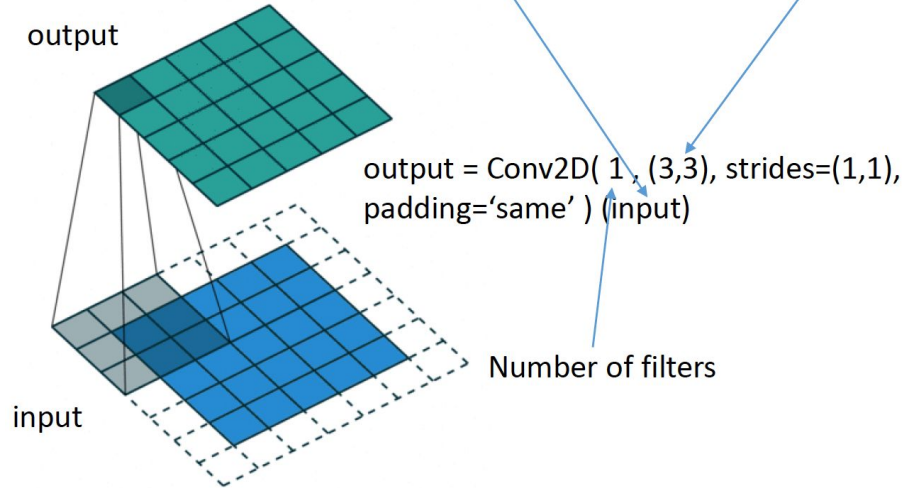
U-Nets explained



https://www.youtube.com/watch?v=NhdzGfB1q74&ab_channel=rupertai

Convolutions + Max Pooling

2D convolution using a kernel size of 3, stride of 1 and padding



Max pooling kernel size of 3, stride of 1, no padding

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

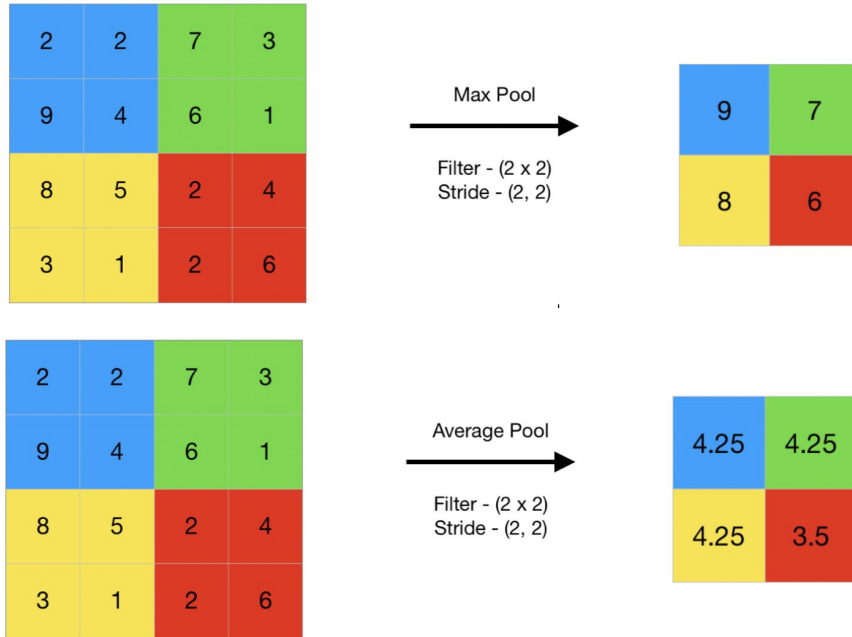
input

3.0	3.0	3.0
3.0	3.0	3.0
3.0	2.0	3.0

output

output = MaxPooling2D((3, 3), strides=(1,1), padding='valid') (input)

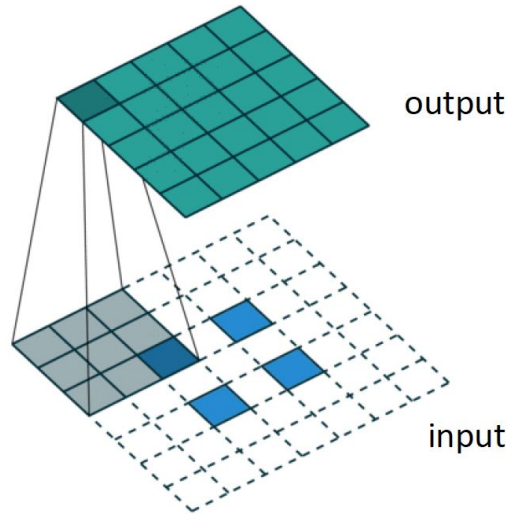
Different pooling layers



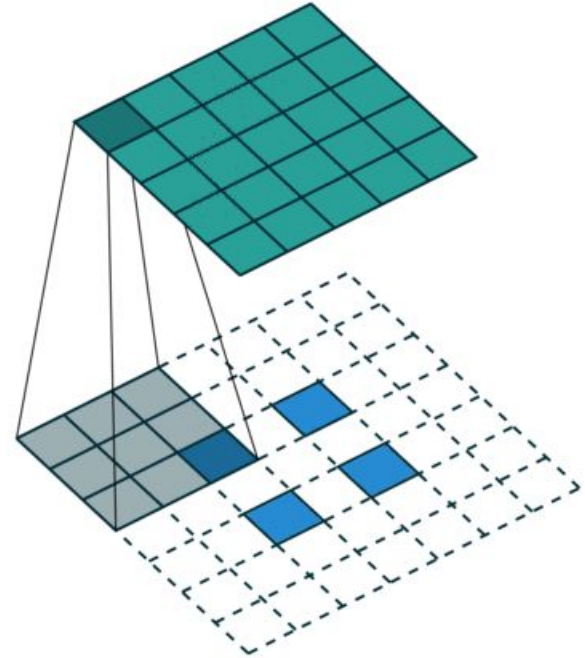
```
1 import numpy as np
2 from keras.models import Sequential
3 from keras.layers import MaxPooling2D
4
5 # define input image
6 image = np.array([[2.0, 2.0, 7.0, 3.0],
7                   [9.0, 4.0, 6.0, 1.0],
8                   [8.0, 5.0, 2.0, 4.0],
9                   [3.0, 1.0, 2.0, 6.0]])
10 image = image.reshape(1.0, 4.0, 4.0, 1.0)
11
12 # define model containing just a single max pooling
13 # layer
14 model = Sequential(
15     [MaxPooling2D(pool_size = 2, strides = 2)])
16
17 # generate pooled output
18 output = model.predict(image)
19
20 # print output image
21 output = np.squeeze(output)
22 print(output)
```

Transpose convolution - Upsampling

Transposed 2D convolution with padding, stride of 2 and kernel of 3



`output = Conv2DTranspose(1, (3,3) , strides= (2,2),
padding='same') (input)`



Transpose convolutions explained by Andrew Ng!

Transpose Convolution

2	1
3	2

2x2

1 ⁻¹	2 ⁻¹	1 ⁻¹
2 ⁻¹	0 ⁻¹	1 ⁻¹
0 ⁻¹	2 ⁻¹	1 ⁻¹



4x4

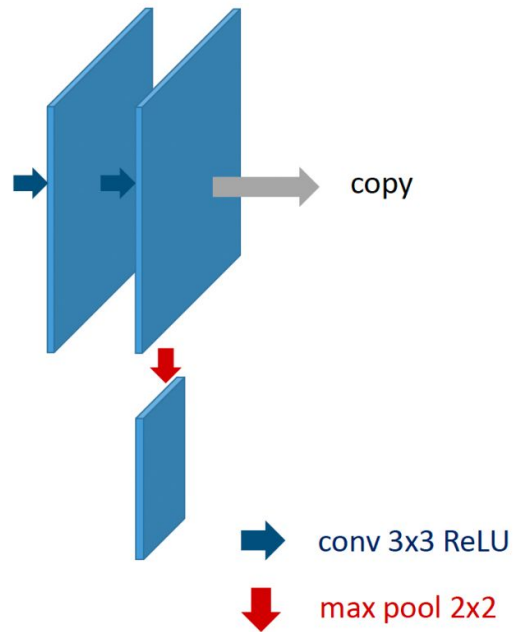
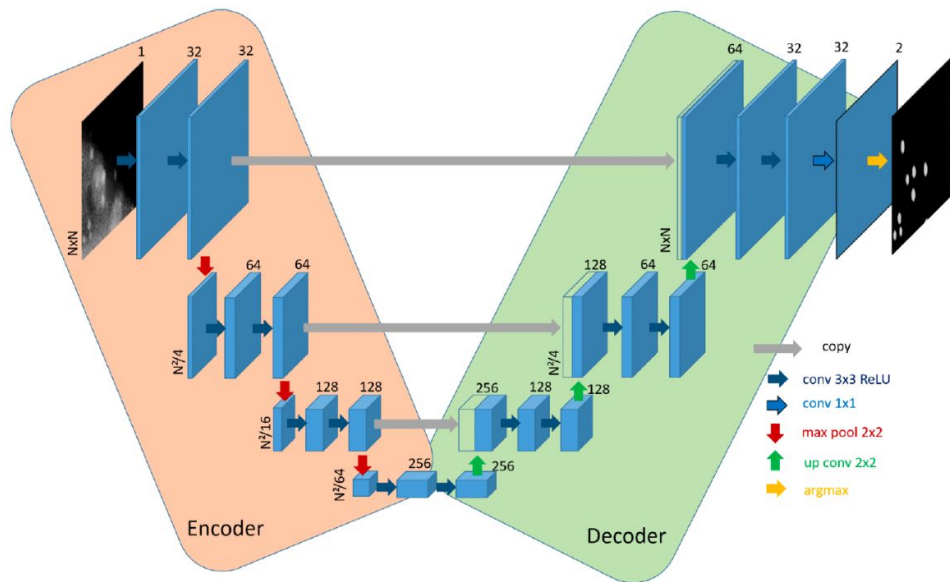
filter $f \times f = 3 \times 3$

padding $p = 1$

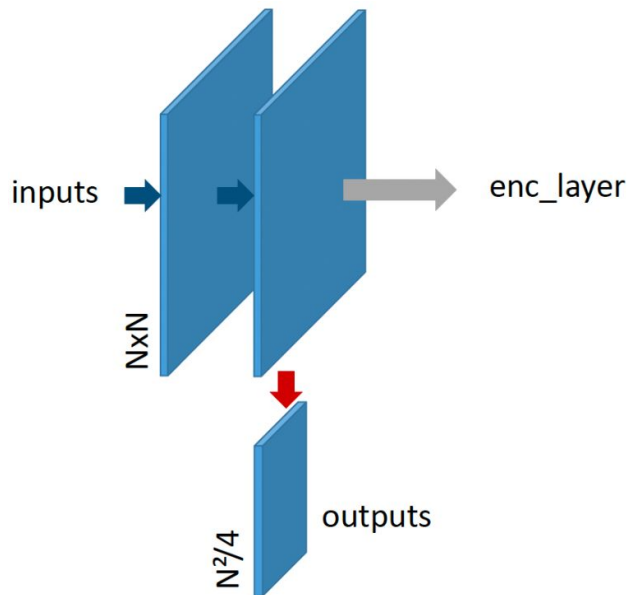
stride $s = 2$

Andrew Ng

UNet: Encoder



UNet: Encoder



➡ conv 3x3 ReLU

⬇ max pool 2x2

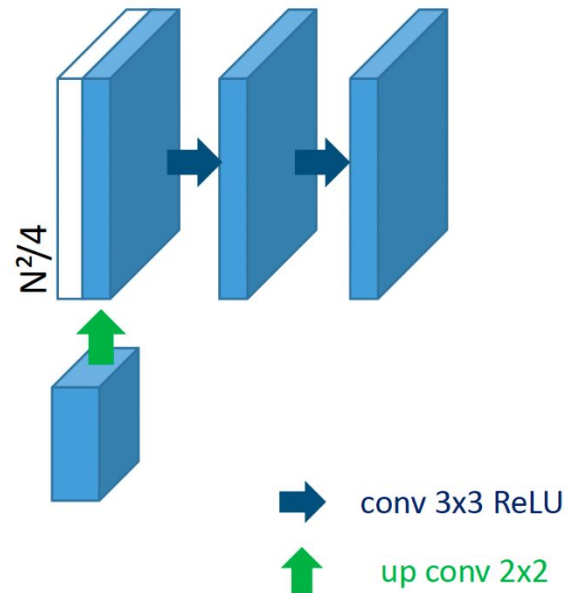
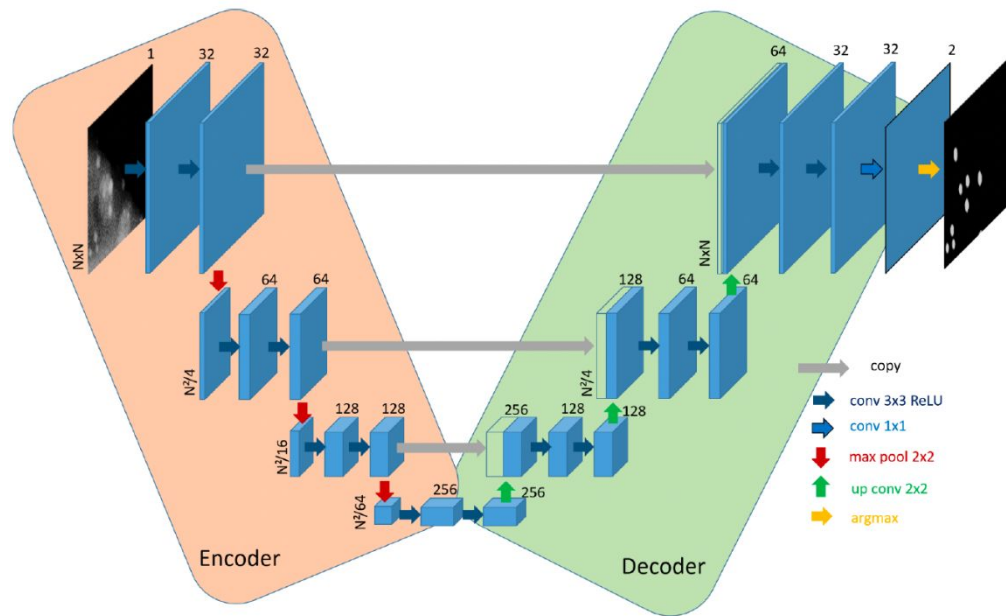
```
c = Conv2D(filters, (3,3), activation='relu',  
kernel_initializer=kernel_initializer, padding='same') (inputs)
```

```
c = Conv2D(filters, (3,3), activation='relu',  
kernel_initializer=kernel_initializer, padding='same') (c)
```

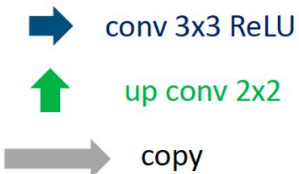
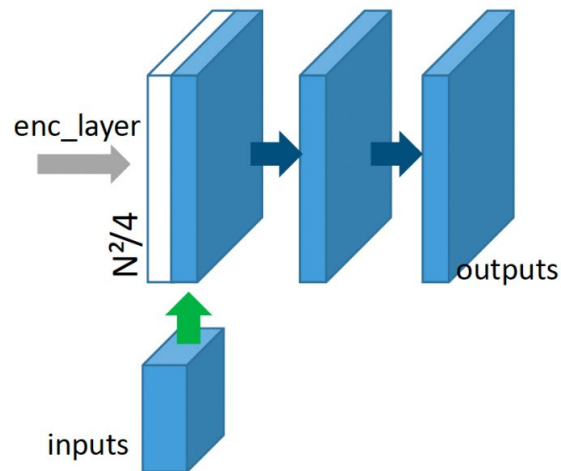
```
enc_layer = c
```

```
outputs= MaxPooling2D((2, 2)) (c)
```

UNet: Decoder



UNet: Decoder



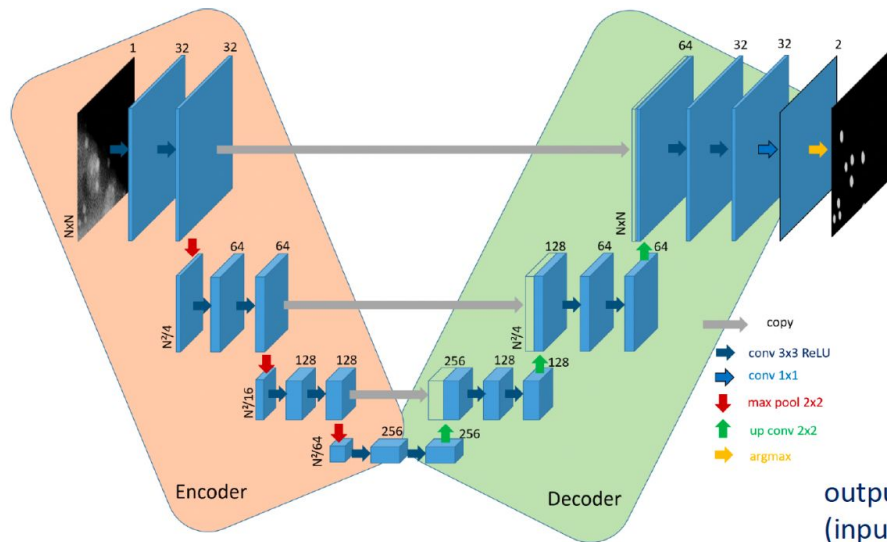
```
c = Conv2DTranspose( filters, (2, 2), strides=(2, 2), padding='same') (input)
```

```
c = Concatenate()([c, enc_layer])
```

```
c = Conv2D(filters, (3,3), activation='relu', kernel_initializer=kernel_initializer,  
padding='same') (inputs)
```

```
outputs = Conv2D(filters, (3,3), activation='relu',  
kernel_initializer=kernel_initializer, padding='same') (c)
```

UNet: Decoder



outputs = Conv2D(num_classes, (1, 1), activation='sigmoid')(input)

The argmax is done outside the network (loss or metric, display, ...)

